

Artificial Intelligence Capstone Project1

Music Genre Classification

110550101- 陳威達

Report Outline

- Introduction
 - About MGC problem
 - About this project
- Data acquisition
 - Dataset
- Feature Extract and Preprocessing
 - VGGish
 - Preprocessing
- Model
 - Model selection and parameter setting
 - Evaluating method
- Result
- Reference

*The detail codes are available on github [\[01\]](#)
and also simply attached as appendix on the last few pages*

Introduction

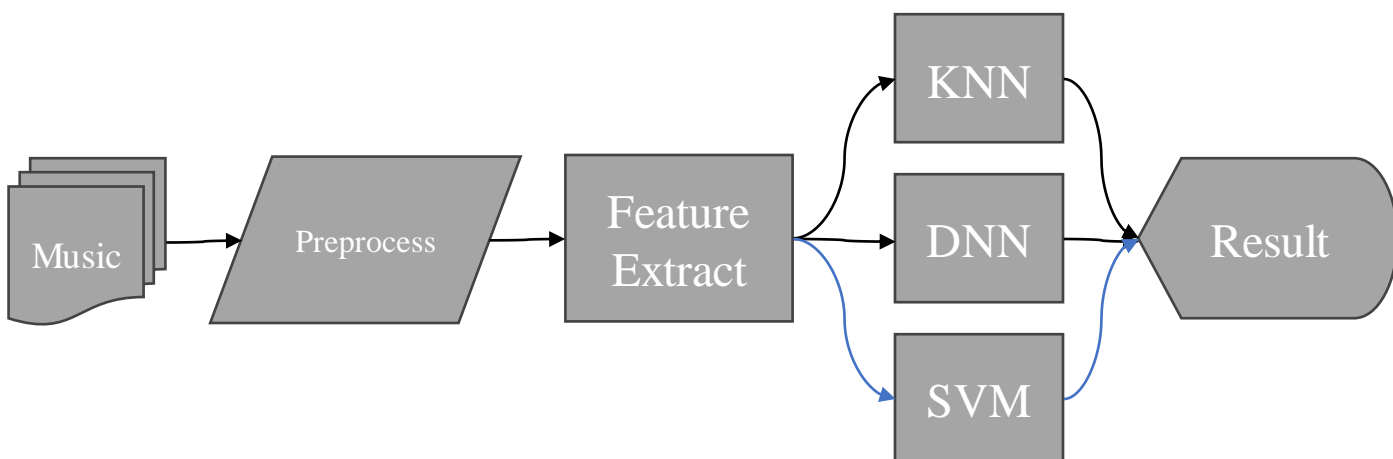
- About MGC

The MGC (Music Genre Classification) problem revolves around the automated categorization of musical compositions using computational methods, a critical task within the realm of music information retrieval. The process involves developing algorithms and techniques dedicated to organizing, searching, and retrieving music data based on distinct musical and cultural characteristics. Music genres serve as pivotal categories utilized by music streaming services, radio stations, festivals, and record labels to group similar music and offer personalized recommendations to users.

The difficult part in MGC problem is that music category that human defines is subjective, which makes the kind of problem inherently challenging. Human perceptions of music genres often carry subjective nuances, influenced by cultural, emotional, and contextual factors. These subjective interpretations introduce complexity as there may be considerable variation in how different individuals categorize and perceive music genres.

- About this project

In this work, I get the inspired by the paper[\[1\]](#) and try to utilize the open source music on YouTube to construct a model to identify the 6 different genres of music. The abstract of model is like:



Data Acquisition

- Dataset

The dataset acquisition method is referred from GTZAN Dataset[2]

The music source is from YouTube studio music library, which is free and generally be allowed to be utilized. Moreover, the library has labeled music genre so I use these labels as the target and download each type of music in this work.

I select six types of music; they are classic, electronic dance music, hip hop , jazz and blue, pop, rock. In each type, I download about 45~50 files and every music file is more than 30 seconds

```
# Get the # of the file in each folder
for music_genre in os.listdir("input"):
    print(music_genre, "contains", len(os.listdir(f"input/{music_genre}")), "files")

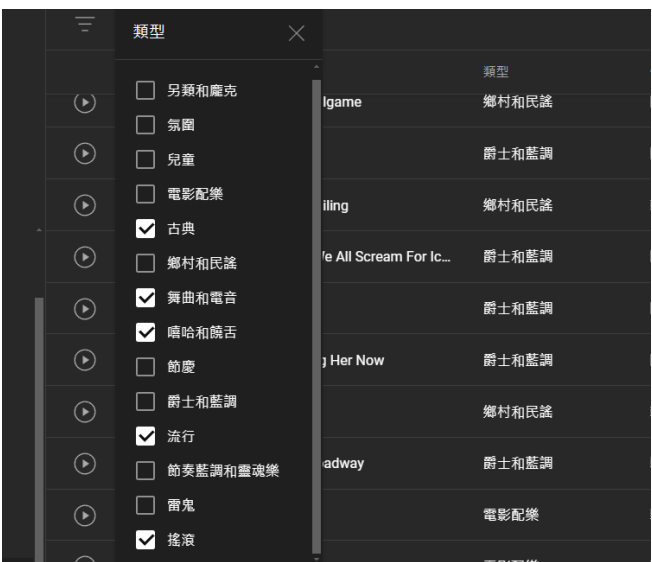
[90] ✓ 0.0s

... classic contains 45 files
edm contains 47 files
hip hop contains 50 files
jazz and blue contains 50 files
pop contains 48 files
rock contains 49 files
```

The amount of music files in each genre

名稱	專輯	位元速率	大小
Healing - Kevin Ma...	YouTube Audio Libr...	320kbps	20,318 KB
I Need to Start Wri...	YouTube Audio Libr...	320kbps	16,865 KB
Mesmerize - Kevin ...	YouTube Audio Libr...	320kbps	16,566 KB
There's Probably N...	YouTube Audio Libr...	320kbps	12,192 KB
John Stockton Slo...	YouTube Audio Libr...	320kbps	11,075 KB
Angevin - Thatche...	YouTube Audio Libr...	320kbps	10,962 KB
Night Snow - Ashe...	YouTube Audio Libr...	320kbps	10,770 KB
Enchanted Journey...	YouTube Audio Libr...	320kbps	10,532 KB
Love of All - Twin ...	YouTube Audio Libr...	320kbps	10,506 KB
The Temperature o...	YouTube Audio Libr...	320kbps	10,321 KB
That Kid in Fourth ...	YouTube Audio Libr...	320kbps	10,289 KB
Edward - Audiona...	YouTube Audio Libr...	320kbps	10,161 KB
Clouds - Huma-Hu...	YouTube Audio Libr...	320kbps	9,556 KB
Life in Romance - T...	YouTube Audio Libr...	320kbps	9,289 KB
Simple Sonata - Sir...	YouTube Audio Libr...	320kbps	9,032 KB
Laserdisc - Chris Z...	YouTube Audio Libr...	320kbps	8,769 KB
Facile - Kevin MacL...	YouTube Audio Libr...	320kbps	8,286 KB
Pastoral - Asher Fu...	YouTube Audio Libr...	320kbps	8,251 KB
Pachabelly - Huma...	YouTube Audio Libr...	320kbps	8,233 KB
Gymnopedie no1 - ...	YouTube Audio Libr...	320kbps	7,508 KB
NirvanaVEVO - Ch...	YouTube Audio Libr...	320kbps	7,471 KB
Enchanted Valley - ...	YouTube Audio Libr...	320kbps	7,435 KB
Lifting Dreams - A...	YouTube Audio Libr...	320kbps	7,346 KB
Gymnopedie No 1 ...	YouTube Audio Libr...	320kbps	7,311 KB
Shattered Paths - ...	YouTube Audio Libr...	320kbps	6,957 KB
The Rain - Silent P...	YouTube Audio Libr...	320kbps	6,874 KB

The appearance of the audio files in the folder



The interface in YouTube Studio music library

Feature Extract and Preprocessing

- VGGish [3]

VGGish is an audio feature extraction model developed by Google that is specifically designed for the task of audio classification, including tasks like environmental sound recognition and music genre classification. It is pre-trained on a large dataset of audio samples and extracts fixed-size embeddings or feature vectors from input audio clips.

In this project, the output of each file is (1,128), which represents the feature of the specific file.

```

Data preprocessing and feature extraction

# Feature extraction
vggish = hub.load('https://tfhub.dev/google/vggish/1')

def vggish_extract(audiofile):
    y, sr = librosa.load(audiofile, sr = 44100)
    window = 20000
    stride = 5000
    total_time = librosa.get_duration(y = y, sr = sr)
    start = 0
    end = total_time * 1000
    return_list = []
    for i in range(start, int(end), stride):
        if i + window > end:
            break
        y_temp = y[i:i+window]
        feature = vggish(y_temp).numpy()

        if feature.shape[0] == 0:
            continue
        return_list.append(feature)
    return return_list

[3] ✓ 0.6s
```

The code for data preprocessing and feature extraction

- Preprocessing

According to the previous study[4], regardless of a smaller number but longer period as the input, the larger quantity of splits with shorter period of audio duration period can extract more accurate features. Hence, I select a 20s window with a 5s stride as range and review each audio file to extract feature. Since each genre has different numbers of files and each file has different period, the numbers of output in each class of music are different.

Model

- Model selection and parameter setting

The requirement in this project is two supervised machine learning model and one unsupervised learning model. Due to the well work of feature selection by VGGish, I don't select complex or pretrained model. In this project, I use a SVM, a fully connected neural network(DNN) and KNN as the algorithm to train my model

In the part of hyper parameters in SVM and KNN, I consider the setting of the study of [\[4\]](#). As for the DNN, I use relatively shallower layers to avoid weight vanishing

```
models = {
    'knn': KNeighborsClassifier(n_neighbors = 1, algorithm= 'brute'),
    'svm': SVC(kernel= 'poly', degree= 6,tol= 0.001, coef0= 0.1 ,gamma= 'scale')
}
```

✓ 0.0s

```
NN = Sequential()
NN.add(Dense(128, input_dim=x.shape[1], activation='relu'))
NN.add(Dense(64, activation='relu'))
NN.add(Dense(class_num, activation='softmax'))
NN.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

✓ 0.0s

```
models["DNN"] = NN
print(NN.summary())
```

✓ 0.0s

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 128)	16512
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 6)	390

=====

Total params: 25,158
Trainable params: 25,158
Non-trainable params: 0

The code for model setting

Model

- Evaluating method

Due to relatively small dataset , I select 10-fold cross validation method to measure the score of accuracy, f1, precision, recall and mcc to get the aggregate performance of each algorithm.

```
# model evaluation
def evaluate_model(predictions, y_test):
    ✨ accuracy = accuracy_score(y_test, predictions)
    f1 = f1_score(y_test, predictions, average='weighted')
    precision = precision_score(y_test, predictions, zero_division=1, average='weighted')
    recall = recall_score(y_test, predictions, average='weighted')
    mcc = matthews_corrcoef(y_test, predictions)
    # auROC = roc_auc_score(y_test, predictions, multi_class= 'ovo')

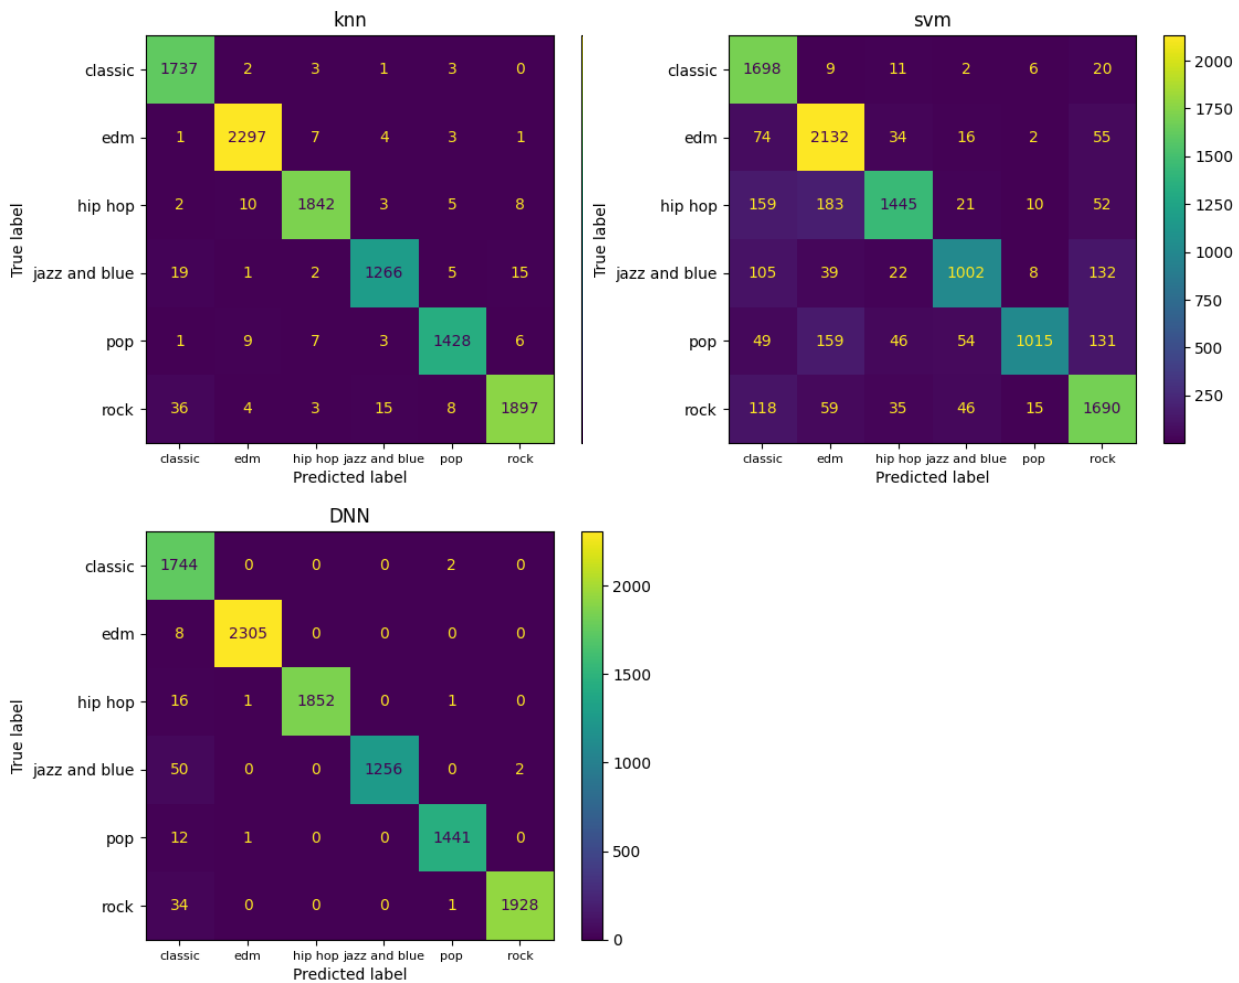
    return {'accuracy': accuracy, 'f1': f1, 'precision': precision, 'recall': recall, 'mcc': mcc}
```

```
kf = KFold(n_splits= 10, shuffle=True, random_state=42)
result = []
for kf_idx, (train_idx, val_idx) in enumerate(kf.split(x)):
    x_train, x_val = x.iloc[train_idx], x.iloc[val_idx]
    y_train, y_val = y[train_idx], y[val_idx]
    for model_name, model in models.items():
        if model_name == 'DNN':
            model.fit(x_train, y_train, epochs=30, batch_size=32)
            predictions = np.argmax(model.predict(x_val), axis=1)
        else:
            model.fit(x_train, y_train)
            predictions = model.predict(x_val)
        fold_res = evaluate_model(predictions, y_val)
        fold_res['model'] = model_name
        fold_res['fold'] = kf_idx
        result.append(fold_res)
```

The code for estimating model performance

Result

In the result presentation, I additionally calculate the mean of previous 10-fold performance and draft the confusion matrix.



```
print("Cross-validation results: ")
print(all_result_df[-3:])
all_result_df.to_csv(output_folder + '/result.csv', index = False)
```

✓ 0.0s

Cross-validation results:

	model	accuracy	f1	precision	recall	mcc
30	DNN Average	0.940976	0.940789	0.942469	0.940976	0.928937
31	knn Average	0.861929	0.860928	0.861433	0.861929	0.833093
32	svm Average	0.760467	0.755032	0.766977	0.760467	0.711998

The result figures and statics

Reference

- A thesis submitted in partial fulfilment of the requirements for the degree of Master of Science in Computer Science [\[1\]](#)
- GTZAN Dataset - Music Genre Classification [\[2\]](#)
- Vggish [\[3\]](#)
- Music Genre Classification: A Review of Deep-Learning and Traditional Machine-Learning Approaches [\[4\]](#)

Appendix

```
# Importing necessary libraries
```

```
# Data preprocessing
```

```
✓import pandas as pd
```

```
● import numpy as np
```

```
import os, librosa
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Visualization
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix
```

```
# Model
```

```
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, matthews_corrcoef
```

```
import numpy as np
```

```
import tensorflow_hub as hub
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# from sklearn.linear_model import LogisticRegression # not used
```

```
from sklearn.svm import SVC
```

```
from keras.models import Sequential
```

```
from keras.layers import Dense, Dropout
```

```
from sklearn.model_selection import KFold
```

```
✓ 9.8s
```

~ Data preprocessing and feature extraction

[+ 程式碼](#)[+ Markdown](#)

```
# Feature extraction
vggish = hub.load('https://tfhub.dev/google/vggish/1')

def vggish_extract(audiofile):
    y, sr = librosa.load(audiofile, sr = 44100)
    window = 20000
    stride = 5000
    total_time = librosa.get_duration(y = y, sr = sr)
    start = 0
    end = total_time * 1000
    return_list = []
    for i in range(start, int(end), stride):
        if i + window > end:
            break
        y_temp = y[i:i+window]
        feature = vggish(y_temp).numpy()

        if feature.shape[0] == 0:
            continue
        return_list.append(feature)
    return return_list
```

[3] ✓ 29.3s

```
# get the current working directory

folder_path = 'input'
class_num = len(os.listdir(folder_path))
data = []
cuurent_amount = 0
for genre in os.listdir(folder_path):
    for file in os.listdir(folder_path + '/' + genre):
        audio_file = folder_path + '/' + genre + '/' + file
        return_list = vggish_extract(audio_file)

        for i in return_list:
            for feature in i.tolist():
                feature.append(genre)
                data.append(feature)

        # print('File: {} is done'.format(file))
# print the ammount of data extracted in each genre
print('Genre: {} is done'.format(genre))
print('The ammount of data extracted is {}'.format(len(data) - cuurent_amount))
cuurent_amount = len(data)
print('_____')
```

```
# divide the data into features and labels
x = data.iloc[:, :-1]
y = data.iloc[:, -1]

# encode the labels
encoder = LabelEncoder()
y = encoder.fit_transform(y)
print(f"data shape x: {x.shape} y: {y.shape}")
```



```
data shape x: (10654, 128) y: (10654,)
```

```
# model evaluation
def evaluate_model(predictions, y_test):

    accuracy = accuracy_score(y_test, predictions)
    f1 = f1_score(y_test, predictions, average='weighted')
    precision = precision_score(y_test, predictions, zero_division=1, average='weighted')
    recall = recall_score(y_test, predictions, average='weighted')
    mcc = matthews_corrcoef(y_test, predictions)
    # auroc = roc_auc_score(y_test, predictions, multi_class= 'ovo')

    return {'accuracy': accuracy, 'f1': f1, 'precision': precision, 'recall': recall, 'mcc': mcc}
```

Traditional model

```
models = {
    'knn': KNeighborsClassifier(n_neighbors = 1, algorithm= 'brute'),
    'svm': SVC(kernel= 'poly', degree= 6,tol= 0.001, coef0= 0.1 ,gamma= 'scale')
}
```

```
[ ]
```

Neural network model

```
NN = Sequential()
NN.add(Dense(128, input_dim=x.shape[1], activation='relu'))
NN.add(Dense(64, activation='relu'))
NN.add(Dense(class_num, activation='softmax'))
NN.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
[ ]
```

```

kf = KFold(n_splits= 10, shuffle=True, random_state=42)
result = []
for kf_idx, (train_idx, val_idx) in enumerate(kf.split(x)):
    x_train, x_val = x.iloc[train_idx], x.iloc[val_idx]
    y_train, y_val = y[train_idx], y[val_idx]
    for model_name, model in models.items():

        if model_name == 'DNN':
            model.fit(x_train, y_train, epochs=30, batch_size=32)
            predictions = np.argmax(model.predict(x_val), axis=1)
        else:
            model.fit(x_train, y_train)
            predictions = model.predict(x_val)
        fold_res = evaluate_model(predictions, y_val)
        fold_res['model'] = model_name
        fold_res['fold'] = kf_idx
        result.append(fold_res)

```

Result presentation

```

output_folder = 'result'
if not os.path.exists(output_folder):
    os.makedirs(output_folder)
else:
    # delete the existing files
    for file in os.listdir(output_folder):
        os.remove(output_folder + '/' + file)

cv_result_df = pd.DataFrame(result)
avg_result = cv_result_df.groupby('model').mean().reset_index()
avg_result['model'] += ' Average '
all_result_df = pd.concat([cv_result_df, avg_result], ignore_index=True)
all_result_df = all_result_df[['model', 'accuracy', 'f1', 'precision', 'recall', 'mcc']]

print("Confusion matrix for the each model: ")
for model_name, model in models.items():
    plt.figure(figsize=(20, 20))
    if model_name == 'DNN':
        predictions = np.argmax(model.predict(x), axis=1)
    else:
        predictions = model.predict(x)
    cm = confusion_matrix(y, predictions)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=encoder.classes_)
    disp.plot()
    plt.title(model_name)
    plt.tick_params(axis='x', labelsz=8)
    plt.savefig(output_folder + '/' + model_name + '.png')
    plt.show()
    print("Confusion matrix for ", model_name, " is saved as ", model_name + '.png')

```